

# MSDS600 Week 5 Assignment Starter - Rafael Fernandes

## Getting Ready


Loading the necessary libraries, dataset, and filters and checking if the data has any error.

```
In [4]: import pandas as pd
        from pycaret.classification import setup, compare_models, predict_model, save_model
        from IPython.display import Code
```

```
In [5]: df = pd.read_csv('churn_data.csv', index_col='customerID')
        df.head(10)
```

Out[5]:

	tenure	PhoneService	Contract	PaymentMethod	MonthlyCharges	TotalCharg
customerID						
7590-VHVEG	1	No	Month-to-month	Electronic check	29.85	29.
5575-GNVDE	34	Yes	One year	Mailed check	56.95	1889.
3668-QPYBK	2	Yes	Month-to-month	Mailed check	53.85	108.
7795-CFOCW	45	No	One year	Bank transfer (automatic)	42.30	1840.
9237-HQITU	2	Yes	Month-to-month	Electronic check	70.70	151.
9305-CDSKC	8	Yes	Month-to-month	Electronic check	99.65	820.
1452-KIOVK	22	Yes	Month-to-month	Credit card (automatic)	89.10	1949.
6713-OKOMC	10	No	Month-to-month	Mailed check	29.75	301.
7892-POOKP	28	Yes	Month-to-month	Electronic check	104.80	3046.
6388-TABGU	62	Yes	One year	Bank transfer (automatic)	56.15	3487.



To use pycaret I created a virtual environment and I called it 'pyca'.

In [7]: `!jupyter kernelspec list`

Available kernels:

```
pyca      C:\Users\rafaf\AppData\Roaming\jupyter\kernels\pyca
python3   C:\Users\rafaf\AppData\Roaming\jupyter\kernels\python3
```

```
0.00s - Debugger warning: It seems that frozen modules are being used, which may
0.01s - make the debugger miss breakpoints. Please pass -Xfrozen_modules=off
0.00s - to python to disable frozen modules.
0.00s - Note: Debugging will proceed. Set PYDEVD_DISABLE_FILE_VALIDATION=1 to disabl
e this validation.
```

In [8]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Index: 7043 entries, 7590-VHVEG to 3186-AJIEK
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   tenure                7043 non-null   int64
1   PhoneService          7043 non-null   object
2   Contract              7043 non-null   object
3   PaymentMethod         7043 non-null   object
4   MonthlyCharges        7043 non-null   float64
5   TotalCharges          7032 non-null   float64
6   Churn                 7043 non-null   object
dtypes: float64(2), int64(1), object(4)
memory usage: 440.2+ KB
```

## Automation

In this part I start the process for auto ML, setting it up, comparing the models and I'm sorting 'recall' as first model.

```
In [10]: automl = ClassificationExperiment()
```

```
In [11]: automl = setup(data=df, target='Churn')
```

	Description	Value
0	Session id	5151
1	Target	Churn
2	Target type	Binary
3	Target mapping	No: 0, Yes: 1
4	Original data shape	(7043, 7)
5	Transformed data shape	(7043, 12)
6	Transformed train set shape	(4930, 12)
7	Transformed test set shape	(2113, 12)
8	Numeric features	3
9	Categorical features	3
10	Rows with missing values	0.2%
11	Preprocess	True
12	Imputation type	simple
13	Numeric imputation	mean
14	Categorical imputation	mode
15	Maximum one-hot encoding	25
16	Encoding method	None
17	Fold Generator	StratifiedKFold
18	Fold Number	10
19	CPU Jobs	-1
20	Use GPU	False
21	Log Experiment	False
22	Experiment Name	clf-default-name
23	USI	2002

```
In [12]: best_model = compare_models(sort='recall')
```

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
<b>lr</b>	Logistic Regression	0.7903	0.8320	0.7903	0.7806	0.7828	0.4279	0.4326	1.1690
<b>gbc</b>	Gradient Boosting Classifier	0.7880	0.8341	0.7880	0.7760	0.7767	0.4068	0.4159	0.1930
<b>lightgbm</b>	Light Gradient Boosting Machine	0.7880	0.8233	0.7880	0.7779	0.7801	0.4201	0.4251	0.1760
<b>ridge</b>	Ridge Classifier	0.7878	0.8221	0.7878	0.7741	0.7746	0.3986	0.4094	0.0420
<b>ada</b>	Ada Boost Classifier	0.7868	0.8336	0.7868	0.7759	0.7777	0.4122	0.4186	0.0970
<b>lda</b>	Linear Discriminant Analysis	0.7856	0.8221	0.7856	0.7761	0.7787	0.4184	0.4220	0.0460
<b>svm</b>	SVM - Linear Kernel	0.7708	0.7426	0.7708	0.7591	0.7527	0.3431	0.3613	0.0490
<b>rf</b>	Random Forest Classifier	0.7602	0.7941	0.7602	0.7494	0.7528	0.3508	0.3539	0.1810
<b>knn</b>	K Neighbors Classifier	0.7560	0.7468	0.7560	0.7413	0.7455	0.3268	0.3316	0.0740
<b>et</b>	Extra Trees Classifier	0.7491	0.7633	0.7491	0.7383	0.7423	0.3244	0.3265	0.2000
<b>dummy</b>	Dummy Classifier	0.7347	0.5000	0.7347	0.5398	0.6223	0.0000	0.0000	0.0580
<b>dt</b>	Decision Tree Classifier	0.7260	0.6530	0.7260	0.7264	0.7260	0.2976	0.2979	0.0480
<b>nb</b>	Naive Bayes	0.6801	0.8065	0.6801	0.7881	0.6987	0.3580	0.4043	0.0530
<b>qda</b>	Quadratic Discriminant Analysis	0.6408	0.6674	0.6408	0.6833	0.6164	0.1475	0.1640	0.0410

Processing: 0%| | 0/61 [00:00<?, ?it/s]

In [13]: automl

Out[13]: <pycaret.classification.oop.ClassificationExperiment at 0x29f53b73210>

In [14]: best\_model


```
Out[14]: LogisticRegression
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                  intercept_scaling=1, l1_ratio=None, max_iter=1000,
                  multi_class='auto', n_jobs=None, penalty='l2',
                  random_state=5151, solver='lbfgs', tol=0.0001, verbose=
0,
                  warm_start=False)
```

Above we can see that 'Logistic Regression' was the best model for the 'Recall' and it showed that 'Accuracy' model with the same result, but AUC had the highest number with 0.8320. Now I'm going to evaluate the model plotting the best model.

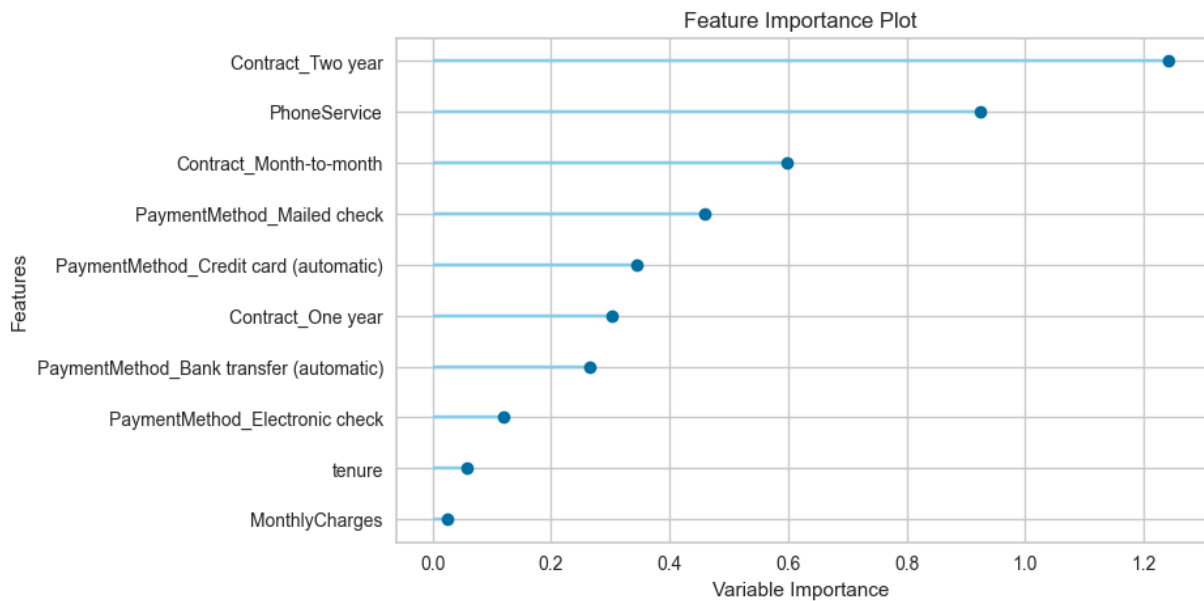
```
In [16]: automl.evaluate_model(best_model)
```

```
interactive(children=(ToggleButtons(description='Plot Type:', icons=('',)), options=
(('Pipeline Plot', 'pipelin...
```

```
In [17]: automl.plot_model(best_model)
```



```
In [18]: automl.plot_model(best_model, plot = 'feature')
```



I'm going to create a new churn data to predict the best model.

```
In [20]: new_churn_data = df.iloc[-3:-2]
```

```
In [21]: predictions = predict_model(best_model, data=new_churn_data)
predictions
```

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
0	Logistic Regression	1.0000	0	1.0000	1.0000	1.0000	nan	0.0000

```
Out[21]:
```

	tenure	PhoneService	Contract	PaymentMethod	MonthlyCharges	TotalCharg
customerID						

4801-JZAZL	11	No	Month-to-month	Electronic check	29.6	346.4500
------------	----	----	----------------	------------------	------	----------

## Saving, testing and loading the model

In this part, I'm going to save the model in a pickle file, then I'm going to test, load, and predict the file.

```
In [23]: automl.save_model(best_model, 'pyca_data_model')
```

Transformation Pipeline and Model Successfully Saved

```

Out[23]: (Pipeline(memory=Memory(location=None),
                steps=[('label_encoding',
                        TransformerWrapperWithInverse(exclude=None, include=None,
                                                         transformer=LabelEncoder()))),
          ('numerical_imputer',
           TransformerWrapper(exclude=None,
                               include=['tenure', 'MonthlyCharges',
                                         'TotalCharges'],
                               transformer=SimpleImputer(add_indicator=False,
                                                           copy=True,
                                                           fill_value=None,
                                                           keep_empty_features
                                                           =False,...
                                                           handle_missing='ret
                                                           urn_nan',
                                                           handle_unknown='val
                                                           ue',
                                                           return_df=True,
                                                           use_cat_names=True,
                                                           verbose=0))),
          ('trained_model',
           LogisticRegression(C=1.0, class_weight=None, dual=False,
                               fit_intercept=True, intercept_scaling=1,
                               l1_ratio=None, max_iter=1000,
                               multi_class='auto', n_jobs=None,
                               penalty='l2', random_state=5151,
                               solver='lbfgs', tol=0.0001, verbose=0,
                               warm_start=False))),
          verbose=False),
          'pyca_data_model.pkl')

```

```

In [24]: pyca_model = ClassificationExperiment()
         tested_model = pyca_model.load_model('pyca_data_model')

```

Transformation Pipeline and Model Successfully Loaded

```

In [25]: new_pyca = ClassificationExperiment()
         loaded_model = new_pyca.load_model('pyca_data_model')

```

Transformation Pipeline and Model Successfully Loaded

```

In [26]: new_pyca.predict_model(loaded_model, df.iloc[-3:-2])

```

Out[26]:                   tenure  PhoneService  Contract  PaymentMethod  MonthlyCharges  TotalCharg

customerID

4801-JZAZL	11	No	Month-to-month	Electronic check	29.6	346.4500
------------	----	----	----------------	------------------	------	----------

◀ ◻ ▶

## Creating a Python module for predictions



```
In [28]: Code('predict_churn.py')
```

```
Out[28]: import pandas as pd
from pycaret.classification import ClassificationExperiment
```

```
def load_data(filepath):
    "Load the churn_data.csv data into a DataFrame."

    df = pd.read_csv('churn_data.csv', index_col='customerID')
    return df

def make_predictions(df):
    "Use the best model (LogisticRegression) pycaret to make predictions"

    classifier = ClassificationExperiment()
    model = classifier.load_model('pyca_data_model')
    predictions = classifier.predict_model(model, data=df)
    predictions.rename({'Label': 'Churn'}, axis=1, inplace=True)
    predictions['Churn'].replace({1: 'Churn', 0: 'No churn'},
                                inplace=True)
    return predictions['Churn']

if __name__ == "__main__":
    df = load_data('churn_data.csv')
    predictions = make_predictions(df)
    print('predictions:')
    print(predictions)
```

Lastly I'm running the file to test it and see the predictions.

```
In [30]: %run predict_churn.py
```

Transformation Pipeline and Model Successfully Loaded

predictions:

customerID

7590-VHVEG No

5575-GNVDE No

3668-QPYBK Yes

7795-CFOCW No

9237-HQITU Yes

...

6840-RESVB No

2234-XADUH No

4801-JZAZL No

8361-LTMKD Yes

3186-AJIEK No

Name: Churn, Length: 7043, dtype: category

Categories (2, object): ['No', 'Yes']

<Figure size 800x550 with 0 Axes>

## References

The following links are references used as resources to complete and improve this project.

[\*\*A step-by-step guide to install PyCaret in Python\*\*](#)

[\*\*A Complete Guide to PyCaret!!!\*\*](#)

[\*\*Analysis and model explainability functions in PyCaret\*\*](#)

[\*\*joblib 1.4.2\*\*](#)

[\*\*FTE\\_Week\\_3 MSDS600 W3 FTE advanced section\*\*](#)

## Summary

I used the pycaret auto ML package to predict if customers are going to churn. I set 'recall' as the metric used for finding the best model and it showed 'Logistic Regression' as the best one, however, 'Accuracy' was the same, and both for all the models had the same result. I trained the model, I plotted the best model and the best model with 'feature'.

After I estimated the predictions for the new DF, I saved the model to the disk as a pickle file, tested the functions with the new data, and printed the predictions Logistic Regression had the best 5 results from 7 comparisons.