

MSDS600 Week 5 Assignment Starter - Rafael Fernandes

Getting Ready


Loading the necessary libraries, dataset, and filters and checking if the data has any error.

```
In [5]: import pandas as pd
        from pycaret.classification import setup, compare_models, predict_model, save_model
        from IPython.display import Code
```

```
In [6]: df = pd.read_csv('churn_data.csv', index_col='customerID')
        df.head(10)
```

Out[6]:

| | tenure | PhoneService | Contract | PaymentMethod | MonthlyCharges | TotalCharg |
|------------|--------|--------------|----------------|---------------------------|----------------|------------|
| customerID | | | | | | |
| 7590-VHVEG | 1 | No | Month-to-month | Electronic check | 29.85 | 29. |
| 5575-GNVDE | 34 | Yes | One year | Mailed check | 56.95 | 1889. |
| 3668-QPYBK | 2 | Yes | Month-to-month | Mailed check | 53.85 | 108. |
| 7795-CFOCW | 45 | No | One year | Bank transfer (automatic) | 42.30 | 1840. |
| 9237-HQITU | 2 | Yes | Month-to-month | Electronic check | 70.70 | 151. |
| 9305-CDSKC | 8 | Yes | Month-to-month | Electronic check | 99.65 | 820. |
| 1452-KIOVK | 22 | Yes | Month-to-month | Credit card (automatic) | 89.10 | 1949. |
| 6713-OKOMC | 10 | No | Month-to-month | Mailed check | 29.75 | 301. |
| 7892-POOKP | 28 | Yes | Month-to-month | Electronic check | 104.80 | 3046. |
| 6388-TABGU | 62 | Yes | One year | Bank transfer (automatic) | 56.15 | 3487. |



To use pycaret I created a virtual environment and I called it 'pyca'.

In [8]: `!jupyter kernelspec list`

Available kernels:

```
pyca      C:\Users\rafaf\AppData\Roaming\jupyter\kernels\pyca
python3   C:\Users\rafaf\AppData\Roaming\jupyter\kernels\python3
```

```
0.01s - Debugger warning: It seems that frozen modules are being used, which may
0.00s - make the debugger miss breakpoints. Please pass -Xfrozen_modules=off
0.00s - to python to disable frozen modules.
0.00s - Note: Debugging will proceed. Set PYDEVD_DISABLE_FILE_VALIDATION=1 to disabl
e this validation.
```

In [9]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Index: 7043 entries, 7590-VHVEG to 3186-AJIEK
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   tenure                7043 non-null   int64
1   PhoneService          7043 non-null   object
2   Contract              7043 non-null   object
3   PaymentMethod         7043 non-null   object
4   MonthlyCharges        7043 non-null   float64
5   TotalCharges          7032 non-null   float64
6   Churn                 7043 non-null   object
dtypes: float64(2), int64(1), object(4)
memory usage: 440.2+ KB
```

Automation

In this part I start the process for auto ML, setting it up, comparing the models and I'm sorting 'recall' as first model.

```
In [11]: automl = ClassificationExperiment()
```

```
In [12]: automl = setup(data=df, target='Churn')
```

| | Description | Value |
|----|-----------------------------|------------------|
| 0 | Session id | 358 |
| 1 | Target | Churn |
| 2 | Target type | Binary |
| 3 | Target mapping | No: 0, Yes: 1 |
| 4 | Original data shape | (7043, 7) |
| 5 | Transformed data shape | (7043, 12) |
| 6 | Transformed train set shape | (4930, 12) |
| 7 | Transformed test set shape | (2113, 12) |
| 8 | Numeric features | 3 |
| 9 | Categorical features | 3 |
| 10 | Rows with missing values | 0.2% |
| 11 | Preprocess | True |
| 12 | Imputation type | simple |
| 13 | Numeric imputation | mean |
| 14 | Categorical imputation | mode |
| 15 | Maximum one-hot encoding | 25 |
| 16 | Encoding method | None |
| 17 | Fold Generator | StratifiedKFold |
| 18 | Fold Number | 10 |
| 19 | CPU Jobs | -1 |
| 20 | Use GPU | False |
| 21 | Log Experiment | False |
| 22 | Experiment Name | clf-default-name |
| 23 | USI | 9673 |

```
In [13]: best_model = compare_models(sort='recall')
```

| | Model | Accuracy | AUC | Recall | Prec. | F1 | Kappa | MCC | TT (Sec) |
|-----------------|---------------------------------|----------|--------|--------|--------|--------|--------|--------|----------|
| lr | Logistic Regression | 0.7931 | 0.8368 | 0.7931 | 0.7834 | 0.7854 | 0.4341 | 0.4392 | 0.9410 |
| gbc | Gradient Boosting Classifier | 0.7931 | 0.8381 | 0.7931 | 0.7835 | 0.7853 | 0.4338 | 0.4392 | 0.1910 |
| ada | Ada Boost Classifier | 0.7917 | 0.8375 | 0.7917 | 0.7812 | 0.7831 | 0.4268 | 0.4327 | 0.1080 |
| lda | Linear Discriminant Analysis | 0.7909 | 0.8265 | 0.7909 | 0.7827 | 0.7848 | 0.4355 | 0.4390 | 0.0390 |
| ridge | Ridge Classifier | 0.7892 | 0.8265 | 0.7892 | 0.7763 | 0.7768 | 0.4051 | 0.4155 | 0.0400 |
| lightgbm | Light Gradient Boosting Machine | 0.7880 | 0.8289 | 0.7880 | 0.7790 | 0.7811 | 0.4248 | 0.4290 | 0.1390 |
| rf | Random Forest Classifier | 0.7661 | 0.7921 | 0.7661 | 0.7573 | 0.7603 | 0.3727 | 0.3750 | 0.1870 |
| knn | K Neighbors Classifier | 0.7611 | 0.7423 | 0.7611 | 0.7466 | 0.7505 | 0.3399 | 0.3449 | 0.0560 |
| et | Extra Trees Classifier | 0.7489 | 0.7690 | 0.7489 | 0.7422 | 0.7446 | 0.3357 | 0.3372 | 0.1550 |
| dummy | Dummy Classifier | 0.7347 | 0.5000 | 0.7347 | 0.5398 | 0.6223 | 0.0000 | 0.0000 | 0.0390 |
| dt | Decision Tree Classifier | 0.7187 | 0.6464 | 0.7187 | 0.7216 | 0.7198 | 0.2851 | 0.2854 | 0.0470 |
| nb | Naive Bayes | 0.6909 | 0.8132 | 0.6909 | 0.7941 | 0.7088 | 0.3760 | 0.4210 | 0.0400 |
| svm | SVM - Linear Kernel | 0.6651 | 0.6823 | 0.6651 | 0.7572 | 0.6476 | 0.2489 | 0.3007 | 0.0470 |
| qda | Quadratic Discriminant Analysis | 0.5751 | 0.6107 | 0.5751 | 0.6854 | 0.5611 | 0.1233 | 0.1523 | 0.0460 |

Processing: 0%| | 0/61 [00:00<?, ?it/s]

In [14]: automl

Out[14]: <pycaret.classification.oop.ClassificationExperiment at 0x2234bef2410>

In [15]: best_model

Out[15]:

```
LogisticRegression
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=1000,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=358, solver='lbfgs', tol=0.0001, verbose=
0,
                    warm_start=False)
```


Above we can see that 'Logistic Regression' was the best model for the 'Recall' and it showed that 'Accuracy' model with the same result.

Now I'm going to evaluate the model plotting the best model.


In [17]: `automl.evaluate_model(best_model)`

```
interactive(children=(ToggleButtons(description='Plot Type:', icons=('',)), options=
(('Pipeline Plot', 'pipelin...
```

In [18]: `automl.plot_model(best_model)`



In [19]: `automl.plot_model(best_model, plot = 'feature')`



I'm going to create a new churn data to predict the best model.

```
In [21]: new_churn_data = df.iloc[-3:-2]
```

```
In [22]: predictions = predict_model(best_model, data=new_churn_data)
predictions
```

| | Model | Accuracy | AUC | Recall | Prec. | F1 | Kappa | MCC |
|---|---------------------|----------|-----|--------|--------|--------|-------|--------|
| 0 | Logistic Regression | 1.0000 | 0 | 1.0000 | 1.0000 | 1.0000 | nan | 0.0000 |

```
Out[22]:
```

| | tenure | PhoneService | Contract | PaymentMethod | MonthlyCharges | TotalCharg |
|------------|--------|--------------|----------------|------------------|----------------|------------|
| customerID | | | | | | |
| 4801-JZAZL | 11 | No | Month-to-month | Electronic check | 29.6 | 346.4500 |

Saving, testing and loading the model

In this part, I'm going to save the model in a pickle file, then I'm going to test, load, and predict the file.

```
In [24]: automl.save_model(best_model, 'pyca_data_model')
```

Transformation Pipeline and Model Successfully Saved

```

Out[24]: (Pipeline(memory=Memory(location=None),
                steps=[('label_encoding',
                        TransformerWrapperWithInverse(exclude=None, include=None,
                                                         transformer=LabelEncoder()))),
          ('numerical_imputer',
           TransformerWrapper(exclude=None,
                               include=['tenure', 'MonthlyCharges',
                                         'TotalCharges'],
                               transformer=SimpleImputer(add_indicator=False,
                                                           copy=True,
                                                           fill_value=None,
                                                           keep_empty_features
                                                           =False, ...
                                                           handle_missing='return_nan',
                                                           handle_unknown='value',
                                                           return_df=True,
                                                           use_cat_names=True,
                                                           verbose=0))),
          ('trained_model',
           LogisticRegression(C=1.0, class_weight=None, dual=False,
                               fit_intercept=True, intercept_scaling=1,
                               l1_ratio=None, max_iter=1000,
                               multi_class='auto', n_jobs=None,
                               penalty='l2', random_state=358,
                               solver='lbfgs', tol=0.0001, verbose=0,
                               warm_start=False))),
          verbose=False),
          'pyca_data_model.pkl')

```

```

In [25]: pyca_model = ClassificationExperiment()
         tested_model = pyca_model.load_model('pyca_data_model')

```

Transformation Pipeline and Model Successfully Loaded

```

In [26]: new_pyca = ClassificationExperiment()
         loaded_model = new_pyca.load_model('pyca_data_model')

```

Transformation Pipeline and Model Successfully Loaded

```

In [27]: new_pyca.predict_model(loaded_model, df.iloc[-3:-2])

```

Out[27]: tenure PhoneService Contract PaymentMethod MonthlyCharges TotalCharg

customerID

| | | | | | | |
|------------|----|----|----------------|------------------|------|----------|
| 4801-JZAZL | 11 | No | Month-to-month | Electronic check | 29.6 | 346.4500 |
|------------|----|----|----------------|------------------|------|----------|



Creating a Python module for predictions


```
In [29]: Code('predict_churn.py')
```

```
Out[29]: import pandas as pd
from pycaret.classification import ClassificationExperiment
```

```
def load_data(filepath):
    "Load the churn_data.csv data into a DataFrame."

    df = pd.read_csv('churn_data.csv', index_col='customerID')
    return df

def make_predictions(df):
    "Use the best model (LogisticRegression) pycaret to make predictions"

    classifier = ClassificationExperiment()
    model = classifier.load_model('pyca_data_model')
    predictions = classifier.predict_model(model, data=df)
    predictions.rename({'Label': 'Churn'}, axis=1, inplace=True)
    predictions['Churn'].replace({1: 'Churn', 0: 'No churn'},
                                inplace=True)
    return predictions['Churn']

if __name__ == "__main__":
    df = load_data('churn_data.csv')
    predictions = make_predictions(df)
    print('predictions:')
    print(predictions)
```

Lastly I'm running the file to test it and see the predictions.

```
In [31]: %run predict_churn.py
```

Transformation Pipeline and Model Successfully Loaded

predictions:

customerID

7590-VHVEG No

5575-GNVDE No

3668-QPYBK Yes

7795-CFOCW No

9237-HQITU Yes

...

6840-RESVB No

2234-XADUH No

4801-JZAZL No

8361-LTMKD Yes

3186-AJIEK No

Name: Churn, Length: 7043, dtype: category

Categories (2, object): ['No', 'Yes']

<Figure size 800x550 with 0 Axes>

References

The following links are references used as resources to complete and improve this project.

[A step-by-step guide to install PyCaret in Python](#)

[A Complete Guide to PyCaret!!!](#)

[Analysis and model explainability functions in PyCaret](#)

[joblib 1.4.2](#)

[FTE_Week_3 MSDS600 W3 FTE advanced section](#)

Summary

I used the pycaret auto ML package to predict if customers are going to churn. I set 'recall' as the metric used for finding the best model and it showed 'Logistic Regression' as the best one, however, 'Accuracy' was the same, and both for all the models had the same result. I trained the model, I plotted the best model and the best model with 'feature'.

After I estimated the predictions for the new DF, I saved the model to the disk as a pickle file, tested the functions with the new data, and printed the predictions.